

# Linear Classification Models

Fabio A. González Ph.D.

Depto. de Ing. de Sistemas e Industrial  
Universidad Nacional de Colombia, Bogotá

April 9, 2018

# Content

- 1 Introduction
- 2 The perceptron
- 3 Logistic regression
- 4 Logistic regression optimization

# Outline

## 1 Introduction

## 2 The perceptron

## 3 Logistic regression

## 4 Logistic regression optimization

## Classification problems

- $predict(x) = \begin{cases} C_1, & y(x) \geq \text{threshold} \\ C_2, & y(x) < \text{threshold} \end{cases}$ ,  
with  $\text{threshold} = 0$  or  $\text{threshold} = 0.5$ .

## Classification problems

- $\text{predict}(x) = \begin{cases} C_1, & y(x) \geq \text{threshold} \\ C_2, & y(x) < \text{threshold} \end{cases}$ ,  
with  $\text{threshold} = 0$  or  $\text{threshold} = 0.5$ .
- Three ways to address the classification problem:

## Classification problems

- $predict(x) = \begin{cases} C_1, & y(x) \geq \text{threshold} \\ C_2, & y(x) < \text{threshold} \end{cases}$ ,  
with  $\text{threshold} = 0$  or  $\text{threshold} = 0.5$ .
- Three ways to address the classification problem:
  - ① Directly model the discrimination function: e.g.  
 $y(x) = w^T x + w_0$

## Classification problems

- $predict(x) = \begin{cases} C_1, & y(x) \geq \text{threshold} \\ C_2, & y(x) < \text{threshold} \end{cases}$ ,  
with  $\text{threshold} = 0$  or  $\text{threshold} = 0.5$ .
- Three ways to address the classification problem:
  - ① Directly model the discrimination function: e.g.  
 $y(x) = w^T x + w_0$
  - ② Generative model:

$$y(x) = P(C_k|x) = \frac{P(x|C_k)P(C_k)}{P(x)}$$

## Classification problems

- $predict(x) = \begin{cases} C_1, & y(x) \geq \text{threshold} \\ C_2, & y(x) < \text{threshold} \end{cases}$ ,  
with  $\text{threshold} = 0$  or  $\text{threshold} = 0.5$ .
- Three ways to address the classification problem:
  - ① Directly model the discrimination function: e.g  
 $y(x) = w^T x + w_0$
  - ② Generative model:

$$y(x) = P(C_k|x) = \frac{P(x|C_k)P(C_k)}{P(x)}$$

- ③ Discriminative model:

$$y(x) = P(C_k|x) = f(x),$$

with  $f$  an arbitrary function.



# Linear Models

$$y(x) = f(w^T x + w_0)$$

- $f(\cdot)$ : activation function, may be non-linear

# Linear Models

$$y(x) = f(w^T x + w_0)$$

- $f(\cdot)$ : activation function, may be non-linear
- Even if  $f(\cdot)$  is non-linear, the decision boundary is linear

# Linear Models

$$y(x) = f(w^T x + w_0)$$

- $f(\cdot)$ : activation function, may be non-linear
- Even if  $f(\cdot)$  is non-linear, the decision boundary is linear
- Also called *generalized linear models*

# Linear Models

$$y(x) = f(w^T x + w_0)$$

- $f(\cdot)$ : activation function, may be non-linear
- Even if  $f(\cdot)$  is non-linear, the decision boundary is linear
- Also called *generalized linear models*
- Applicable if instead of  $x$  we use a vector of basis functions  $\phi(x)$ , corresponding to features in a feature space

## Using regression for classification

- We can use a regression model, such as least squares to fit a linear classification model with a linear activation function

$$\min_{w, w_0} \sum_{i=1}^{\ell} (t_i - w^T x_i + w_0)^2,$$

where  $t_i \in \{-1, 1\}$  is the label of the  $i$ -th training sample,


## Using regression for classification

- We can use a regression model, such as least squares to fit a linear classification model with a linear activation function

$$\min_{w, w_0} \sum_{i=1}^{\ell} (t_i - w^T x_i + w_0)^2,$$

where  $t_i \in \{-1, 1\}$  is the label of the  $i$ -th training sample,

- but this strategy does not work well:



# Outline

- 1 Introduction
- 2 The perceptron**
- 3 Logistic regression
- 4 Logistic regression optimization

# Rosemblatt's perceptron



- Designed by Frank Rossemlat in 1957



# Rosemblatt's perceptron



- Designed by Frank Rossemlat in 1957
- A hardware implementation of the learning algorithm

# Rosemblatt's perceptron



- Designed by Frank Rossemlat in 1957
- A hardware implementation of the learning algorithm
- The precursor of neural networks

## Rosemblatt's perceptron



- Designed by Frank Rossemblat in 1957
- A hardware implementation of the learning algorithm
- The precursor of neural networks
- Criticized by Marvin Minsky, producing a decline in research funding

# Perceptron learning

- Activation function:

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

# Perceptron learning

- Activation function:

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

- Loss function:

$$E_p(w) = - \sum_{n=1}^{\ell} w^T \phi_n t_n,$$

where  $\phi_n = \phi(x_n)$ .

# Perceptron learning

- Activation function:

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

- Loss function:

$$E_p(w) = - \sum_{n=1}^{\ell} w^T \phi_n t_n,$$

where  $\phi_n = \phi(x_n)$ .

- Learning rule:

$$w^{(\tau+1)} = w^{(\tau)} + \eta \phi_n t_n$$

# Perceptron convergence

- If the training points are linearly separable, the perceptron algorithm converges (*perceptron convergence theorem*)

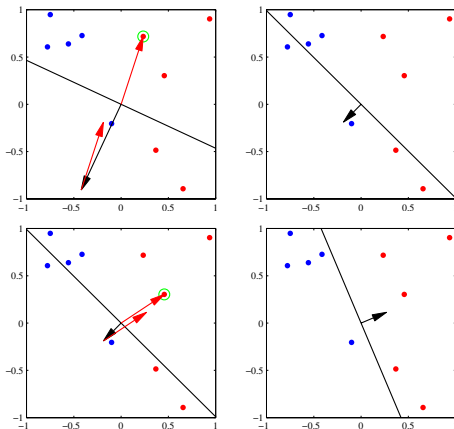
## Perceptron convergence

- If the training points are linearly separable, the perceptron algorithm converges (*perceptron convergence theorem*)
- It could converge to different solutions depending on the order of presentation of training sample



# Perceptron convergence

- If the training points are linearly separable, the perceptron algorithms converges (*perceptron convergence theorem*)
- It could converge to different solutions depending on the order of presentation of training sample



# Perceptron problems

- Non-probabilistic outputs

# Perceptron problems

- Non-probabilistic outputs
- Non-convex optimization problem

# Perceptron problems

- Non-probabilistic outputs
- Non-convex optimization problem
- No convergence guarantee if samples are not linearly separable

# Perceptron problems

- Non-probabilistic outputs
- Non-convex optimization problem
- No convergence guarantee if samples are not linearly separable
- Its power may be increased by stacking several perceptron layers (multilayer perceptrons) and using smooth activation functions

# Outline

- 1 Introduction
- 2 The perceptron
- 3 Logistic regression
- 4 Logistic regression optimization

# Parametric discrimination

- These three conditions are equivalent:
  - $P(C_1|x) \geq 0.5$
  - $\frac{P(C_1|x)}{1-P(C_1|x)} \geq 1$
  - $\text{logit}(P(C_1|x)) = \log \frac{P(C_1|x)}{1-P(C_1|x)} \geq 0$

## Parametric discrimination

- These three conditions are equivalent:
  - $P(C_1|x) \geq 0.5$
  - $\frac{P(C_1|x)}{1-P(C_1|x)} \geq 1$
  - $\text{logit}(P(C_1|x)) = \log \frac{P(C_1|x)}{1-P(C_1|x)} \geq 0$
- If we assume that  $P(x|C_1)$  and  $P(x|C_2)$  are normally distributed sharing the same covariance matrix:

$$\text{logit}(P(C_1|x)) = \log \frac{P(C_1|x)}{P(C_2|x)} = w^T x + w_0,$$

where

$$w = \Sigma^{-1}(\mu_1 - \mu_2)$$

$$w_0 = -\frac{1}{2}(\mu_1 + \mu_2)^T \Sigma^{-1}(\mu_1 + \mu_2) + \log \frac{P(C_1)}{P(C_2)}$$



# Logistic function

- The logit function:

$$\text{logit}(P(C_1|x)) = \log \frac{P(C_1|x)}{1 - P(C_1|x)} = w^T x + w_0$$

## Logistic function

- The logit function:

$$\text{logit}(P(C_1|x)) = \log \frac{P(C_1|x)}{1 - P(C_1|x)} = w^T x + w_0$$

- The inverse-logit:

$$P(C_1|x) = \sigma(w^T x + w_0) = \frac{1}{1 + e^{-(w^T x + w_0)}}$$

# Logistic function


- The logit function:

$$\text{logit}(P(C_1|x)) = \log \frac{P(C_1|x)}{1 - P(C_1|x)} = w^T x + w_0$$

- The inverse-logit:

$$P(C_1|x) = \sigma(w^T x + w_0) = \frac{1}{1 + e^{-(w^T x + w_0)}}$$

- $\sigma$  is called the logistic or sigmoid function.



# Logistic regression



$$y(x) = P(C_1|x) = \sigma(w^T x)$$

# Logistic regression



$$y(x) = P(C_1|x) = \sigma(w^T x)$$

- Find  $w$  using maximum likelihood estimation:

$$p(\mathbf{t}|w) = \prod_{n=1}^{\ell} y_n^{t_n} (1 - y_n)^{1-t_n},$$

where  $\mathbf{t} = \{t_1, \dots, t_\ell\}$  and  $y_n = y(x_n)$ .

# Logistic regression



$$y(x) = P(C_1|x) = \sigma(w^T x)$$

- Find  $w$  using maximum likelihood estimation:

$$p(\mathbf{t}|w) = \prod_{n=1}^{\ell} y_n^{t_n} (1 - y_n)^{1-t_n},$$

where  $\mathbf{t} = \{t_1, \dots, t_\ell\}$  and  $y_n = y(x_n)$ .

- Cross-entropy error:

$$E(w) = -\ln p(\mathbf{t}|w) = -\sum_{n=1}^{\ell} [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)]$$

# Multiclass logistic regression

$$y_k(x) = P(C_k|x) = \frac{e^{w_k^T x}}{\sum_j e^{w_j^T x}}$$

# Multiclass logistic regression



$$y_k(x) = P(C_k|x) = \frac{e^{w_k^T x}}{\sum_j e^{w_j^T x}}$$

- Likelihood:

$$p(\mathbf{T}|w_1 \dots w_K) = \prod_{n=1}^{\ell} \prod_{k=1}^K y_{nk}^{t_{nk}},$$

where  $y_{nk} = y_k(x_n)$  and  $\mathbf{T} \in \mathbb{R}^{\ell \times K}$  is a matrix of target variables with elements  $t_{nk}$ .



# Multiclass logistic regression



$$y_k(x) = P(C_k|x) = \frac{e^{w_k^T x}}{\sum_j e^{w_j^T x}}$$

- Likelihood:

$$p(\mathbf{T}|w_1 \dots w_K) = \prod_{n=1}^{\ell} \prod_{k=1}^K y_{nk}^{t_{nk}},$$

where  $y_{nk} = y_k(x_n)$  and  $\mathbf{T} \in \mathbb{R}^{\ell \times K}$  is a matrix of target variables with elements  $t_{nk}$ .

- Multiclass cross-entropy error:

$$E(w_1, \dots, w_K) = -\ln p(\mathbf{T}|w_1 \dots w_K) = -\sum_{n=1}^{\ell} \sum_{k=1}^K t_{nk} \ln y_{nk}$$

# Outline

- 1 Introduction
- 2 The perceptron
- 3 Logistic regression
- 4 Logistic regression optimization

# Optimization problem



$$\min_w E(w) = \min_w - \sum_{n=1}^{\ell} [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)]$$

# Optimization problem

•

$$\min_w E(w) = \min_w - \sum_{n=1}^{\ell} [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)]$$

•

$$\nabla E(w) = \sum_{n=1}^{\ell} (y_n - t_n) \phi_n$$

# Optimization problem

- $$\min_w E(w) = \min_w - \sum_{n=1}^{\ell} [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)]$$

- $$\nabla E(w) = \sum_{n=1}^{\ell} (y_n - t_n) \phi_n$$

- $$w^{(\tau+1)} = w^{(\tau)} - \eta \sum_{n=1}^{\ell} (y_n - t_n) \phi_n$$

# Newton-Raphson



$$w^{(\tau+1)} = w^{(\tau)} - H^{-1} \nabla E(w)$$

# Newton-Raphson



$$w^{(\tau+1)} = w^{(\tau)} - H^{-1} \nabla E(w)$$



$$\nabla E(w) = \Phi^T(\mathbf{y} - \mathbf{t})$$

# Newton-Raphson



$$w^{(\tau+1)} = w^{(\tau)} - \mathbf{H}^{-1} \nabla E(w)$$



$$\nabla E(w) = \Phi^T (\mathbf{y} - \mathbf{t})$$



$$\mathbf{H} = \nabla \nabla E(w) = \Phi^T \mathbf{R} \Phi,$$

with  $\mathbf{R}$  a diagonal matrix with  $R_{nn} = y_n(1 - y_n)$ .



# Newton-Raphson

- 

$$w^{(\tau+1)} = w^{(\tau)} - \mathbf{H}^{-1} \nabla E(w)$$

- 

$$\nabla E(w) = \Phi^T (\mathbf{y} - \mathbf{t})$$

- 

$$\mathbf{H} = \nabla \nabla E(w) = \Phi^T \mathbf{R} \Phi,$$

with  $\mathbf{R}$  a diagonal matrix with  $R_{nn} = y_n(1 - y_n)$ .

- The resulting algorithm is called *iterative reweighted least squares*.

# Regularization



$$\min_w -C \sum_{n=1}^{\ell} [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)] + \|w\|^2$$

# Regularization

- 

$$\min_w -C \sum_{n=1}^{\ell} [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)] + \|w\|^2$$

- Prevents overfitting.

# Regularization



$$\min_w -C \sum_{n=1}^{\ell} [t_n \ln y_n + (1 - t_n) \ln(1 - y_n)] + \|w\|^2$$

- Prevents overfitting.
- Equivalent to the inclusion of a prior and finding a MAP solution for  $W$ .

# Stochastic gradient descent



$$\min_w Q(w) = \min_w \sum_{i=1}^n Q_i(w)$$

# Stochastic gradient descent



$$\min_w Q(w) = \min_w \sum_{i=1}^n Q_i(w)$$

- Batch gradient descent:

$$w^{(\tau+1)} = w^{(\tau)} - \alpha \nabla Q(w) = w^{(\tau)} - \alpha \sum_{i=1}^n \nabla Q_i(w)$$

# Stochastic gradient descent



$$\min_w Q(w) = \min_w \sum_{i=1}^n Q_i(w)$$

- Batch gradient descent:

$$w^{(\tau+1)} = w^{(\tau)} - \alpha \nabla Q(w) = w^{(\tau)} - \alpha \sum_{i=1}^n \nabla Q_i(w)$$

- on-line gradient descent:

$$w^{(\tau+1)} = w^{(\tau)} - \alpha \nabla Q_i(w)$$



Alpaydin, E. 2010 Introduction to Machine Learning, 2Ed.  
The MIT Press. (Chap 10)



Russell, S and Norvig, P. 2010 Artificial Intelligence: a  
Modern Approach, 3rd Ed, Prentice-Hall (Sect 18.6)